

CSE534 HOMEWORK 3:

PART A

A1) We will be creating the below topology on mininet using the mininet Python API. It consists of 2 hosts, H1 and H2, connected through a network of 4 routers R1-R4. The python file that creates this topology can be found within the *partA/MyTopo.py* path. H1 is connected to R1 in subnet 10.1.1.0/24, and H2 is connected to R4 in subnet 10.1.6.0/24. The result of the pingAll command can be found too, which indicates that all the nodes in the network are connected to each other as there are no message drops. Note, all experiments were conducted on an EC2 instance running Ubuntu version 20.0.

```
#
#                                     TOPOLOGY
#
#                                     (r2-eth0)  (r2-eth1)
#                                     10.1.2.1/24 | 10.1.4.1/24
#                                     (r1-eth1) 10.1.2.2/24 +-----R2-----+ 10.1.4.2/24 (r4-eth1)
#                                     /
#                                     \
# 10.1.1.2/24 | 10.1.1.1/24 /
#                                     \ 10.1.6.1/24 | 10.1.6.2/24
# H1-----R1
# (h1-eth0) (r1-eth0) \
#                                     / (r4-eth0) (h2-eth0)
#                                     /
# 10.1.3.2/24 +-----R3-----+ 10.1.5.2/24
# (r1-eth2) 10.1.3.1/24 | 10.1.5.1/24 (r4-eth2)
#                                     (r3-eth0) (r3-eth1)
```

a) The network topology

```
Pinging all nodes and routers
*** Ping: testing ping reachability
H1 -> H2 R1 R2 R3 R4
H2 -> H1 R1 R2 R3 R4
R1 -> H1 H2 R2 R3 R4
R2 -> H1 H2 R1 R3 R4
R3 -> H1 H2 R1 R2 R4
R4 -> H1 H2 R1 R2 R3
*** Results: 0% dropped (30/30 received)
```

b) Mininet pingAll command output for static routing

A2)

a) Please find the below screenshot containing the routing table information on all routers. Static routing has been enabled, and we can see an entry on each router for every subnet. We use the following command to setup static routing for every interface which is not reachable directly:

- `> ip route add <destination_subnet> via <neighbor_subnet> dev <local_interface>`

Routing Table on Router R1:						
Kernel IP routing table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface
10.1.1.0	0.0.0.0	255.255.255.0	U	0	0	0 r1-eth0
10.1.2.0	0.0.0.0	255.255.255.0	U	0	0	0 r1-eth1
10.1.3.0	0.0.0.0	255.255.255.0	U	0	0	0 r1-eth2
10.1.4.0	10.1.2.1	255.255.255.0	UG	0	0	0 r1-eth1
10.1.5.0	10.1.3.1	255.255.255.0	UG	0	0	0 r1-eth2
10.1.6.0	10.1.2.1	255.255.255.0	UG	0	0	0 r1-eth1

Routing Table on Router R2:						
Kernel IP routing table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface
10.1.1.0	10.1.2.2	255.255.255.0	UG	0	0	0 r2-eth0
10.1.2.0	0.0.0.0	255.255.255.0	U	0	0	0 r2-eth0
10.1.3.0	10.1.2.2	255.255.255.0	UG	0	0	0 r2-eth0
10.1.4.0	0.0.0.0	255.255.255.0	U	0	0	0 r2-eth1
10.1.5.0	10.1.4.2	255.255.255.0	UG	0	0	0 r2-eth1
10.1.6.0	10.1.4.2	255.255.255.0	UG	0	0	0 r2-eth1

Routing Table on Router R3:						
Kernel IP routing table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface
10.1.1.0	10.1.3.2	255.255.255.0	UG	0	0	0 r3-eth0
10.1.2.0	10.1.3.2	255.255.255.0	UG	0	0	0 r3-eth0
10.1.3.0	0.0.0.0	255.255.255.0	U	0	0	0 r3-eth0
10.1.4.0	10.1.5.2	255.255.255.0	UG	0	0	0 r3-eth1
10.1.5.0	0.0.0.0	255.255.255.0	U	0	0	0 r3-eth1
10.1.6.0	10.1.5.2	255.255.255.0	UG	0	0	0 r3-eth1

Routing Table on Router R4:						
Kernel IP routing table						
Destination	Gateway	Genmask	Flags	Metric	Ref	Use Iface
10.1.1.0	10.1.5.1	255.255.255.0	UG	0	0	0 r4-eth2
10.1.2.0	10.1.4.1	255.255.255.0	UG	0	0	0 r4-eth1
10.1.3.0	10.1.5.1	255.255.255.0	UG	0	0	0 r4-eth2
10.1.4.0	0.0.0.0	255.255.255.0	U	0	0	0 r4-eth1
10.1.5.0	0.0.0.0	255.255.255.0	U	0	0	0 r4-eth2
10.1.6.0	0.0.0.0	255.255.255.0	U	0	0	0 r4-eth0

b) The trace output between H1 and H2

```
mininet> H1 traceroute H2
traceroute to 10.1.6.2 (10.1.6.2), 30 hops max, 60 byte packets
 1  10.1.1.1 (10.1.1.1)  0.032 ms  0.007 ms  0.006 ms
 2  10.1.2.1 (10.1.2.1)  0.016 ms  0.010 ms  0.009 ms
 3  10.1.4.2 (10.1.4.2)  0.019 ms  0.012 ms  0.019 ms
 4  10.1.6.2 (10.1.6.2)  0.021 ms  0.018 ms  0.015 ms
mininet> H2 traceroute H1
traceroute to 10.1.1.2 (10.1.1.2), 30 hops max, 60 byte packets
 1  10.1.6.1 (10.1.6.1)  0.033 ms  0.009 ms  0.008 ms
 2  10.1.4.1 (10.1.4.1)  0.017 ms  0.011 ms  0.010 ms
 3  10.1.2.2 (10.1.2.2)  0.020 ms  0.012 ms  0.013 ms
 4  10.1.1.2 (10.1.1.2)  0.023 ms  0.015 ms  0.015 ms
mininet> 
```

PART B

B1) Please find the below screenshot containing the routing table information on all routers and hosts. We need to spawn a BIRD daemon on every router to enable dynamic routing using the RIP protocol. The main configuration file, *bird.conf* can be found for every router in its corresponding folder (R1/R2/R3/R4) under the main partB subdirectory in the project root directory. The python file, *myRIP.py*, can be found there as well. Also below, the pingAll output.

```

Routing Table on Router R1:
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
10.1.1.0       0.0.0.0         255.255.255.0   U        0      0      0 r1-eth0
10.1.1.0       0.0.0.0         255.255.255.0   U       32      0      0 r1-eth0
10.1.2.0       0.0.0.0         255.255.255.0   U        0      0      0 r1-eth1
10.1.2.0       0.0.0.0         255.255.255.0   U       32      0      0 r1-eth1
10.1.3.0       0.0.0.0         255.255.255.0   U        0      0      0 r1-eth2
10.1.3.0       0.0.0.0         255.255.255.0   U       32      0      0 r1-eth2
10.1.4.0       10.1.2.1        255.255.255.0   UG       32      0      0 r1-eth1
10.1.5.0       10.1.3.1        255.255.255.0   UG       32      0      0 r1-eth2

Routing Table on Router R2:
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
10.1.1.0       10.1.2.2        255.255.255.0   UG       32      0      0 r2-eth0
10.1.2.0       0.0.0.0         255.255.255.0   U        0      0      0 r2-eth0
10.1.2.0       0.0.0.0         255.255.255.0   U       32      0      0 r2-eth0
10.1.3.0       10.1.2.2        255.255.255.0   UG       32      0      0 r2-eth0
10.1.4.0       0.0.0.0         255.255.255.0   U        0      0      0 r2-eth1
10.1.4.0       0.0.0.0         255.255.255.0   U       32      0      0 r2-eth1
10.1.5.0       10.1.4.2        255.255.255.0   UG       32      0      0 r2-eth1
10.1.6.0       10.1.4.2        255.255.255.0   UG       32      0      0 r2-eth1

Routing Table on Router R3:
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
10.1.1.0       10.1.3.2        255.255.255.0   UG       32      0      0 r3-eth0
10.1.2.0       10.1.3.2        255.255.255.0   UG       32      0      0 r3-eth0
10.1.3.0       0.0.0.0         255.255.255.0   U        0      0      0 r3-eth0
10.1.3.0       0.0.0.0         255.255.255.0   U       32      0      0 r3-eth0
10.1.4.0       10.1.5.2        255.255.255.0   UG       32      0      0 r3-eth1
10.1.5.0       0.0.0.0         255.255.255.0   U        0      0      0 r3-eth1
10.1.5.0       0.0.0.0         255.255.255.0   U       32      0      0 r3-eth1
10.1.6.0       10.1.5.2        255.255.255.0   UG       32      0      0 r3-eth1

Routing Table on Router R4:
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
10.1.1.0       10.1.4.1        255.255.255.0   UG       32      0      0 r4-eth1
10.1.2.0       10.1.4.1        255.255.255.0   UG       32      0      0 r4-eth1
10.1.3.0       10.1.5.1        255.255.255.0   UG       32      0      0 r4-eth2
10.1.4.0       0.0.0.0         255.255.255.0   U        0      0      0 r4-eth1
10.1.4.0       0.0.0.0         255.255.255.0   U       32      0      0 r4-eth1
10.1.5.0       0.0.0.0         255.255.255.0   U        0      0      0 r4-eth2
10.1.5.0       0.0.0.0         255.255.255.0   U       32      0      0 r4-eth2
10.1.6.0       0.0.0.0         255.255.255.0   U        0      0      0 r4-eth0
10.1.6.0       0.0.0.0         255.255.255.0   U       32      0      0 r4-eth0

Pinging all nodes and routers
*** Ping: testing ping reachability
H1 -> H2 R1 R2 R3 R4
H2 -> H1 R1 R2 R3 R4
R1 -> H1 H2 R2 R3 R4
R2 -> H1 H2 R1 R3 R4
R3 -> H1 H2 R1 R2 R4
R4 -> H1 H2 R1 R2 R3
*** Results: 0% dropped (30/30 received)
0.0*** Starting CLI:
mininet> H1 route
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
default        10.1.1.1        0.0.0.0         UG        0      0      0 h1-eth0
10.1.1.0       0.0.0.0         255.255.255.0   U         0      0      0 h1-eth0
mininet> H2 route
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
default        10.1.6.1        0.0.0.0         UG        0      0      0 h2-eth0
10.1.6.0       0.0.0.0         255.255.255.0   U         0      0      0 h2-eth0
mininet>

```

We can also see the traceroute output between hosts H1 and H2 below, with all links connected.

```
mininet> H1 traceroute H2
traceroute to 10.1.6.2 (10.1.6.2), 30 hops max, 60 byte packets
 1  10.1.1.1 (10.1.1.1)  0.035 ms  0.009 ms  0.006 ms
 2  10.1.2.1 (10.1.2.1)  0.020 ms  0.021 ms  0.010 ms
 3  10.1.4.2 (10.1.4.2)  0.020 ms  0.013 ms  0.012 ms
 4  10.1.6.2 (10.1.6.2)  0.022 ms  0.015 ms  0.016 ms
mininet> H2 traceroute H1
traceroute to 10.1.1.2 (10.1.1.2), 30 hops max, 60 byte packets
 1  10.1.6.1 (10.1.6.1)  0.034 ms  0.008 ms  0.006 ms
 2  10.1.4.1 (10.1.4.1)  0.016 ms  0.010 ms  0.008 ms
 3  10.1.2.2 (10.1.2.2)  0.019 ms  0.011 ms  0.011 ms
 4  10.1.1.2 (10.1.1.2)  0.020 ms  0.014 ms  0.013 ms
mininet> █
```

B2) We then bring down the **R1-R2** link with the following command:

- `> link R1 R2 down`

We can now see the updated traceroute output as shown below. The new route goes through the 10.1.3.x/24 and 10.1.5.x/24 subnets (previously 10.1.2.x/24 and 10.1.4.x/24).

```
mininet> link R1 R2 down
mininet> H1 traceroute H2
traceroute to 10.1.6.2 (10.1.6.2), 30 hops max, 60 byte packets
 1  10.1.1.1 (10.1.1.1)  0.038 ms  0.009 ms  0.006 ms
 2  10.1.3.1 (10.1.3.1)  0.020 ms  0.010 ms  0.009 ms
 3  10.1.5.2 (10.1.5.2)  0.022 ms  0.013 ms  0.012 ms
 4  10.1.6.2 (10.1.6.2)  0.022 ms  0.016 ms  0.014 ms
mininet> H2 traceroute H1
traceroute to 10.1.1.2 (10.1.1.2), 30 hops max, 60 byte packets
 1  10.1.6.1 (10.1.6.1)  0.030 ms  0.007 ms  0.007 ms
 2  10.1.5.1 (10.1.5.1)  0.017 ms  0.024 ms  0.009 ms
 3  10.1.3.2 (10.1.3.2)  0.019 ms  0.028 ms  0.012 ms
 4  10.1.1.2 (10.1.1.2)  0.021 ms  0.015 ms  0.014 ms
mininet> █
```

PART C

In this part, we use the below 2 commands to limit router bandwidth + buffer size and introduce delays in every interface of all routers with the `tc tbf` and `tc netem` options respectively.

- `> tc qdisc add dev <interface> root handle 1: tbf rate <bandwidth> burst <burst_rate> limit <buffer_size>`
- `> tc qdisc add dev <interface> parent 1:1 handle 10: netem delay <delay_in_ms>`

BIRD is used to run the RIP protocol for dynamic routing. Once the routers have been reconfigured with the above commands, we then spawn xterm instances of hosts H1 and H2 to run the network performance tool, `Iperf`, which will simulate a TCP performance test between the two hosts, with H2 acting as the `Iperf` server and H1 the client. We do 3 measurements, with the router buffer size varying as 10Kb, 5Mb and 25Mb across them. Delay = 30ms and Bandwidth = 100Mbps are kept constant in all 3 runs.

To run H2 as the `Iperf` server, we use the below command on its xterm terminal

- `> iperf3 -s`

To run H1 as the `Iperf` client, we use the below command on its xterm terminal

- `> iperf3 -c <H2_IP>`

We can also add the `-J` option and pipe the result of both server and client to create output json files.

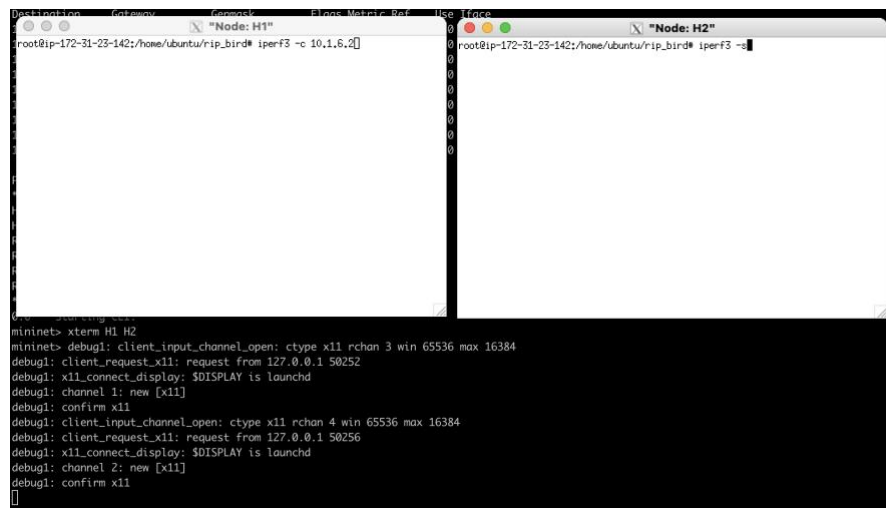
We first run the *MyIperf.py* file in the partC directory that automates the configuration of the routers with the *tc* commands. We can then verify that the routers have been configured correctly using the below 2 commands:

```
mininet> R1 tc qdisc show
qdisc noqueue 0: dev lo root refcnt 2
qdisc tbfb 1: dev r1-eth0 root refcnt 2 rate 100Mbit burst 25Mb lat 0us
qdisc netem 10: dev r1-eth0 parent 1:1 limit 1000 delay 30.0ms
qdisc tbfb 1: dev r1-eth1 root refcnt 2 rate 100Mbit burst 25Mb lat 0us
qdisc netem 10: dev r1-eth1 parent 1:1 limit 1000 delay 30.0ms
qdisc tbfb 1: dev r1-eth2 root refcnt 2 rate 100Mbit burst 25Mb lat 0us
qdisc netem 10: dev r1-eth2 parent 1:1 limit 1000 delay 30.0ms
mininet> H1 ping -c 5 H2
PING 10.1.6.2 (10.1.6.2) 56(84) bytes of data:
64 bytes from 10.1.6.2: icmp_seq=1 ttl=61 time=180 ms
64 bytes from 10.1.6.2: icmp_seq=2 ttl=61 time=180 ms
64 bytes from 10.1.6.2: icmp_seq=3 ttl=61 time=180 ms
64 bytes from 10.1.6.2: icmp_seq=4 ttl=61 time=180 ms
64 bytes from 10.1.6.2: icmp_seq=5 ttl=61 time=180 ms

--- 10.1.6.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4005ms
rtt min/avg/max/mdev = 180.364/180.424/180.476/0.041 ms
mininet>
```

From the above image, we can see that all interfaces of router R1 have been correctly configured by the *tc* command to update its bandwidth, buffer_size and delay parameters. Similarly, we can verify the same for the other routers as well. Also from the ping command, we see that the RTT delay is now 180ms. Every packet flows through 6 interfaces from H1 to H2 (refer topology diagram in Part A, the interfaces are H1 -> r1-eth0 -> r1-eth1 -> r2-eth0 -> r2-eth1 -> r4-eth1 -> r4-eth0 -> H2), and hence total delay is $6 * 30 = 180\text{ms}$ per packet.

Now that we have setup the routers, we will run *iperf* between H1 and H2 through their xterms, as shown below. We will then describe the results for buffer_size = 10Kb, 5Mb and 25Mb.



```
mininet> xterm H1 H2
mininet> debug: client_input_channel_open: ctype x11 rchan 3 win 65536 max 16384
debug: client_request_x11: request from 127.0.0.1 50252
debug: x11_connect_display: $DISPLAY is launchd
debug: channel 1: new [x11]
debug: confirm x11
debug: client_input_channel_open: ctype x11 rchan 4 win 65536 max 16384
debug: client_request_x11: request from 127.0.0.1 50256
debug: x11_connect_display: $DISPLAY is launchd
debug: channel 2: new [x11]
debug: confirm x11
```

We also capture the client/server result json files for each buffer_size configuration and store them in the *partC/iperf_results* directory.



```
root@ip-172-31-23-142:/home/ubuntu/rip_bird# iperf3 -s -J > server_10kb.json
root@ip-172-31-23-142:/home/ubuntu/rip_bird# iperf3 -c 10.1.6.2 -J > client_10kb.json
root@ip-172-31-23-142:/home/ubuntu/rip_bird# ls | grep 10kb
client_10kb.json
server_10kb.json
root@ip-172-31-23-142:/home/ubuntu/rip_bird#
```

We can then estimate the actual bandwidth, number of retransmissions and the round-trip times of the entire flow from the 'bits_per_second', 'retransmits' and 'mean_rtt' attributes (in `$.end.streams[0].sender.* JSONPath`) in the `client_*.json` files, as shown below. The **Bandwidth Delay Product (BDP)** can then be calculated as $\text{BDP} = \text{bits_per_second} * \text{mean_rtt}$

```

"end": {
  "streams": [
    {
      "sender": {
        "socket": 7,
        "start": 0,
        "end": 10.000179,
        "seconds": 10.000179,
        "bytes": 112317088,
        "bits_per_second": 89852062.04808934,
        "retransmits": 5,
        "max_snd_cwnd": 2829392,
        "max_rtt": 271434,
        "min_rtt": 180893,
        "mean_rtt": 215799,
        "sender": true
      },
      "receiver": {
        "socket": 7,
        "start": 0,
        "end": 10.234367,
        "seconds": 10.000179,
        "bytes": 103481320,
        "bits_per_second": 80889278.252382383,
        "sender": true
      }
    }
  ]
}

```

Test 1 – Using buffer_size = 10Kb (Refer to client_10k.json file):

Node: H2

Server listening on 5201

Accepted connection from 10.1.1.2, port 54730

[7]	local 10.1.6.2 port 5201 connected to 10.1.1.2 port 54732		
[ID]	Interval	Transfer	Bitrate
[7]	0.00-1.00 sec	181 KBytes	1.48 Mb/s
[7]	1.00-2.00 sec	5.10 MBytes	42.7 Mb/s
[7]	2.00-3.00 sec	10.9 MBytes	91.8 Mb/s
[7]	3.00-4.00 sec	11.4 MBytes	95.4 Mb/s
[7]	4.00-5.00 sec	11.4 MBytes	95.6 Mb/s
[7]	5.00-6.00 sec	11.4 MBytes	95.6 Mb/s
[7]	6.00-7.00 sec	11.4 MBytes	95.6 Mb/s
[7]	7.00-8.00 sec	11.4 MBytes	95.6 Mb/s
[7]	8.00-9.00 sec	11.4 MBytes	95.6 Mb/s
[7]	9.00-10.00 sec	11.4 MBytes	95.6 Mb/s
[7]	10.00-10.24 sec	2.68 MBytes	95.6 Mb/s

[ID]	Interval	Transfer	Bitrate
[7]	0.00-10.24 sec	98.7 MBytes	80.9 Mb/s

receiver

Server listening on 5201

Node: H1

root@ip-172-31-23-142:/home/ubuntu/rip_bird# iperf3 -c 10.1.6.2

Connecting to host 10.1.6.2, port 5201

[7]	local 10.1.1.2 port 54732 connected to 10.1.6.2 port 5201				
[ID]	Interval	Transfer	Bitrate	Retr	Cwnd
[7]	0.00-1.00 sec	872 KBytes	7.15 Mb/s	0	195 KBytes
[7]	1.00-2.00 sec	15.9 MBytes	133 Mb/s	0	2.70 MBytes
[7]	2.00-3.00 sec	10.0 MBytes	83.9 Mb/s	7	2.05 MBytes
[7]	3.00-4.00 sec	11.2 MBytes	94.4 Mb/s	0	2.22 MBytes
[7]	4.00-5.00 sec	12.5 MBytes	105 Mb/s	0	2.35 MBytes
[7]	5.00-6.00 sec	11.2 MBytes	94.4 Mb/s	0	2.46 MBytes
[7]	6.00-7.00 sec	11.2 MBytes	94.4 Mb/s	0	2.55 MBytes
[7]	7.00-8.00 sec	11.2 MBytes	94.4 Mb/s	0	2.61 MBytes
[7]	8.00-9.00 sec	11.2 MBytes	94.4 Mb/s	0	2.66 MBytes
[7]	9.00-10.00 sec	11.2 MBytes	94.4 Mb/s	0	2.69 MBytes

[ID]	Interval	Transfer	Bitrate	Retr
[7]	0.00-10.00 sec	107 MBytes	89.5 Mb/s	7
[7]	0.00-10.24 sec	98.7 MBytes	80.9 Mb/s	

sender

receiver

iperf Done.

root@ip-172-31-23-142:/home/ubuntu/rip_bird#

a) $\text{BDP} = 89852062.04808934 * 215799 * 10^{-6} = 2.42 \text{ Mbytes}$ (comparable to avg cwnd in H1 above)

Test 2 – Using buffer_size = 5Mb (Refer to client_5m.json file):

Node: H2

Server listening on 5201

Accepted connection from 10.1.1.2, port 54776

[7]	local 10.1.6.2 port 5201 connected to 10.1.1.2 port 54778		
[ID]	Interval	Transfer	Bitrate
[7]	0.00-1.00 sec	167 KBytes	1.37 Mb/s
[7]	1.00-2.00 sec	9.33 MBytes	78.2 Mb/s
[7]	2.00-3.00 sec	11.6 MBytes	97.6 Mb/s
[7]	3.00-4.00 sec	11.4 MBytes	95.9 Mb/s
[7]	4.00-5.00 sec	11.4 MBytes	95.6 Mb/s
[7]	5.00-6.00 sec	11.4 MBytes	95.7 Mb/s
[7]	6.00-7.00 sec	11.4 MBytes	95.6 Mb/s
[7]	7.00-8.00 sec	11.4 MBytes	95.7 Mb/s
[7]	8.00-9.00 sec	11.4 MBytes	95.6 Mb/s
[7]	9.00-10.00 sec	11.4 MBytes	95.5 Mb/s
[7]	10.00-10.70 sec	7.97 MBytes	95.5 Mb/s

[ID]	Interval	Transfer	Bitrate
[7]	0.00-10.70 sec	109 MBytes	85.4 Mb/s

receiver

Server listening on 5201

Node: H1

root@ip-172-31-23-142:/home/ubuntu/rip_bird# iperf3 -c 10.1.6.2

Connecting to host 10.1.6.2, port 5201

[7]	local 10.1.1.2 port 54778 connected to 10.1.6.2 port 5201				
[ID]	Interval	Transfer	Bitrate	Retr	Cwnd
[7]	0.00-1.00 sec	820 KBytes	6.72 Mb/s	0	177 KBytes
[7]	1.00-2.00 sec	19.7 MBytes	166 Mb/s	0	9.50 MBytes
[7]	2.00-3.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	3.00-4.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	4.00-5.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	5.00-6.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	6.00-7.00 sec	12.5 MBytes	105 Mb/s	0	16.0 MBytes
[7]	7.00-8.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	8.00-9.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes
[7]	9.00-10.00 sec	11.2 MBytes	94.4 Mb/s	0	16.0 MBytes

[ID]	Interval	Transfer	Bitrate	Retr
[7]	0.00-10.00 sec	112 MBytes	93.8 Mb/s	0
[7]	0.00-10.70 sec	109 MBytes	85.4 Mb/s	0

sender receiver

iperf Done.

root@ip-172-31-23-142:/home/ubuntu/rip_bird#

a) $\text{BDP} = 93934247.55341047 * 596035 * 10^{-6} = 7 \text{ Mbytes}$

Test 3 – Using buffer size = 25Mb (Refer to client 25m.json file):

```

"Node: H2"
Server listening on 5201
Accepted connection from 10.1.1.2, port 54792
[ 7] local 10.1.6.2 port 5201 connected to 10.1.1.2 port 54794
[ ID] Interval      Transfer    Bitrate
[ 7] 0.00-1.00 sec    181 KBytes  1.48 Mbits/sec
[ 7] 1.00-2.00 sec    9.92 MBytes 83.3 Mbits/sec
[ 7] 2.00-3.00 sec   30.5 MBytes 256 Mbits/sec
[ 7] 3.00-4.00 sec   11.4 MBytes 95.4 Mbits/sec
[ 7] 4.00-5.00 sec   11.4 MBytes 95.9 Mbits/sec
[ 7] 5.00-6.00 sec   11.4 MBytes 95.4 Mbits/sec
[ 7] 6.00-7.00 sec   11.4 MBytes 95.9 Mbits/sec
[ 7] 7.00-8.00 sec   11.4 MBytes 95.4 Mbits/sec
[ 7] 8.00-9.00 sec   11.4 MBytes 95.9 Mbits/sec
[ 7] 9.00-10.00 sec  11.4 MBytes 95.4 Mbits/sec
[ 7] 10.00-10.70 sec  7.95 MBytes 95.5 Mbits/sec
[ ID] Interval      Transfer    Bitrate
[ 7] 0.00-10.70 sec  128 MBytes 101 Mbits/sec
Server listening on 5201
receiver

"Node: H1"
root@ip-172-31-23-142:/home/ubuntu/rip_bird# iperf3 -c 10.1.6.2
Connecting to host 10.1.6.2, port 5201
[ 7] local 10.1.1.2 port 54794 connected to 10.1.6.2 port 5201
[ ID] Interval      Transfer    Bitrate    Retr    Cwnd
[ 7] 0.00-1.00 sec    926 KBytes  7.59 Mbits/sec  0    195 KBytes
[ 7] 1.00-2.00 sec   20.0 MBytes 168 Mbits/sec  0   10.1 MBytes
[ 7] 2.00-3.00 sec   31.2 MBytes 262 Mbits/sec  0   16.5 MBytes
[ 7] 3.00-4.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 4.00-5.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 5.00-6.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 6.00-7.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 7.00-8.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 8.00-9.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ 7] 9.00-10.00 sec   11.2 MBytes 94.4 Mbits/sec  0   16.5 MBytes
[ ID] Interval      Transfer    Bitrate    Retr
[ 7] 0.00-10.00 sec  131 MBytes 110 Mbits/sec  0
[ 7] 0.00-10.70 sec  128 MBytes 101 Mbits/sec  0
iperf Done.
root@ip-172-31-23-142:/home/ubuntu/rip_bird#
sender
receiver

```

a) $BDP = 109744553.52140266 * 565538 * 10^{-6} = 7.76 \text{ Mbytes}$

OBERVATIONS OF THE ABOVE MEASUREMENTS:

- We see some retransmissions in the first measurement with 10Kb buffer size. This may be due to packets being dropped from one of the router interfaces due to the small-size buffer queue filling up. For the other 2 measurements, we see that there are no retransmissions after the router buffer size was increased.
- The BDP is increasing with increase in buffer_size. This is because of increase in the actual average bandwidth (refer Bitrate in H1) with increase in buffer_size.
- We see a major multiplicative increase of the cwnd in the first 2 seconds, followed by a sharp drop in the bandwidth (bitrate) which then leads to the cwnd stabilizing across successive intervals.
- Since retransmission does not occur for buffer_sizes 5Mb and 25Mb, the retransmission time is 0. For 10Kb, the retransmission time may be around 0.24s (10.24s – 10s).